

# CSSE2010/CSSE7201

## Lecture 12



✓  
IN - Afmoy 32x7  
EX - Afmoy 32xP  
✓

# Instruction Set Architecture – Part 1

School of Information Technology and Electrical Engineering  
The University of Queensland

# Instruction Set Architecture (ISA)

- **ISA** defines the interface between hardware and software
  - ISA is a specification
  - Microarchitecture is how the control unit is built
- For hardware (microarchitecture) designers
  - Don't need to know about the high level software
  - Just build a microarchitecture that implements the ISA
- For software writers (machine language programmers and compiler writers)
  - Don't need to know (much) about microarchitecture
  - Just write or generate instructions that match the ISA



# What makes an ISA?


- ✓ 1. Memory models ✓
- ✓ 2. Registers ✓
- ✓ 3. Instructions ✓
- ✓ 4. Data types

✓ { Atmega324A  
Atmega328P }

# What makes an ISA?

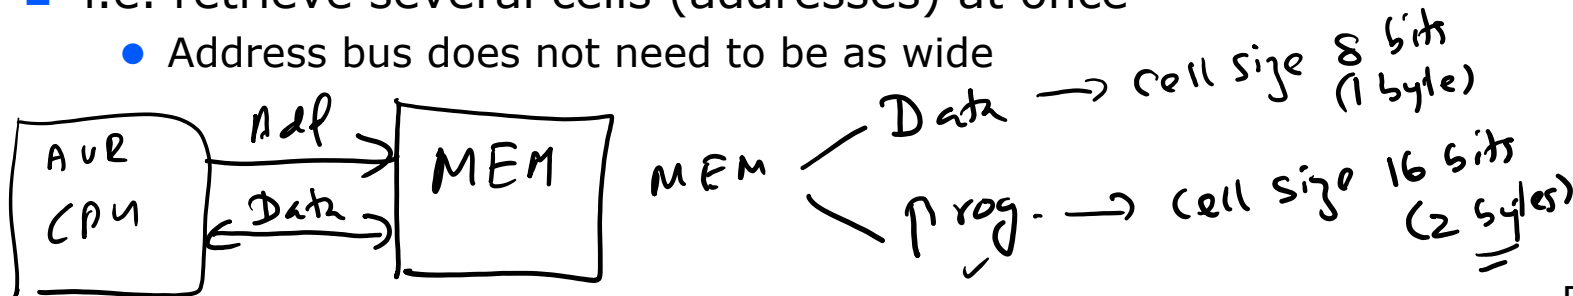
## 1: Memory Models

- ✓ ● Memory model: how does memory look to CPU?
- Issues
  - A. Addressable cell size
  - B. Alignment
  - C. Address spaces
  - D. Endianness




# A. Addressable Cell Size

- Memory has cells, each of which has an unique address (or cell number).
- Most common cell size is 8 bits (1 byte).
  - But not always!
    - AVR Instruction memory has 16 bit cells
- Bus is used to transport the content of cell, but sometimes the data bus may be wider:
  - i.e. retrieve several cells (addresses) at once
    - Address bus does not need to be as wide




# Bus sizes

- For every doubling of data bus width
  - Remove least significant bit of address bus
- Example:  $2^n$  cells x 8-bit cell size
  - Address size: n-bits
- Bus configurations:
  - Data bus: 8 bits      Address bus: n bits
    - One cell transferred at a time
  - Data bus: 16 bits      Address bus: n-1 bits
    - Two cells transferred at a time
  - Data bus: 32 bits      Address bus: n-2 bits
    - Four cells transferred at a time



Data bus width  
Address bus width  
cell size  
# cells



$2^n$  cells  
Address - n bits

If a CPU supports  $2^{35}$  bytes of memory and has a 64 bit data bus, how many bits wide is the address bus?

(cell size is 8 bits)

13% A. 30

13% B. 31

13% C. 32 ✓✓

13% D. 33

13% E. 34

13% F. 35

13% G. 38

13% H. None of the above

Data bus Address bus

8

35

16

34

32

33

64


32 ✓

35 bytes  
2  
cell - 1 byte  
35 cells  
2

# B. Alignment


- Many architectures require **natural alignment**, e.g.
  - 4-byte words starting at addresses 0, 4, 8, ...
  - 8-byte words starting at addresses 0, 8, 16, ...
- For example, with a 8 byte word, it can be stored as follows:

64 bit  
data bus



1 read

Aligned 8-byte word at address 8



2 read

Nonaligned 8-byte word at address 12



## B. Alignment (cont.)

- Alignment often required because it is more efficient
- Example – Pentium IV
  - Fetches 8 bytes at a time from memory (64-bit wide data bus)
  - Alignment is NOT required (so the processor is backwards compatible with earlier processors)
    - e.g. 4-byte word stored at address 4 takes 1 read (bytes 0 to 7).
    - e.g. 4-byte word stored at address 6: Must read bytes 0 to 7 (one read) and bytes 8 to 15 (second read) then extract the 4 required bytes from the 16 bytes read

How many memory reads are needed to read an 8-byte value on a machine with a 16-bit data bus that does not require natural alignment?

0% **A.** Always 2

0% **B.** 2 or 3

0% **C.** Always 3

0% **D.** 3 or 4


0% **E.** Always 4

0% **F.** 4 or 5 ✓✓

0% **G.** Always 5



Case (i)  
4 read operations



Case (ii)  
5 read operations.

# C. Address Spaces

✓

$n$   
 $2^n$

- Many microprocessors have a single linear memory address space (**von Neumann** architecture)

- e.g.  $2^{32}$  bytes numbered from 0 to  $2^{32} - 1$ 
  - (may not be bytes – depends on addressable cell size)

- However, **Harvard architecture** is:

- Separate address spaces for instructions and data
- AVR ATmega 324A**

• Data address space:  $2^{11}$  (2048) bytes ✓

• Instruction address space:  $2^{14}$  16-bit words ✓

– 16,384 instruction words = 32,768 bytes

$32 \cdot 2^{10}$   
 $2^5 \cdot 2^{10}$   
 $2^{15}$   
 $2^{14}$  cells

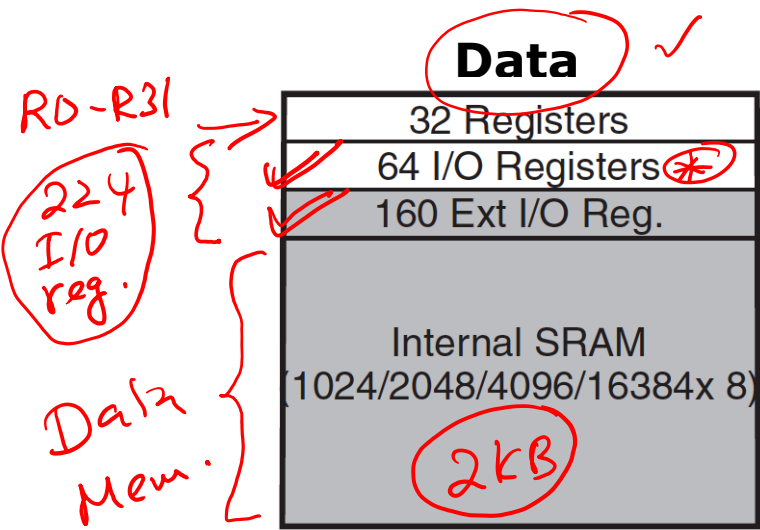
$2 \cdot 2^{10}$   
Data  
2KB  
SRAM  
cell size 8 bits.

✓  
32kB  
Flash

AVR  
CPU

Prog Mem. → cell size (16 bits)


# AVR Address Spaces



← 8 bits →


[From ATmega324A datasheet]

## Program Memory



0x0000

ADD R1, R1



MC


$$2^{14} \cdot 2 = 2^{15} = 2 \cdot 2^{14} = 2 \cdot 2^{10} \cdot 2^4 = 2 \cdot 2^4 \cdot 2^4 = 2 \cdot 16 \cdot 16 = 512 \text{ kB}$$

← 16 bits →

# D. Endianness ✓

- Different machines may support different byte orderings
- **Little endian** – little end (least significant byte) stored first (at lowest address)
  - Intel microprocessors (Pentium etc)
- **Big endian** – big end stored first
  - SPARC, Freescale (formerly Motorola) microprocessors
- Most desktop/workstation/server CPUs produced since ~1992 are "bi-endian" (support both)
  - some switchable at boot time
  - others at run time (i.e. can change dynamically)

Ordering of the bytes  
(not bits)



# Example

- The contents of a memory are shown below. What is the value of:
  - the 1-byte word at address 1? *00110011*
  - the 2-byte word at address 0 if this is a big endian system?

*00001111 00110011*

*Bigendian*

0	00001111
1	00110011 ✓
2	11011101
3	00101100

The contents of a memory are shown below. What is the value of the 2-byte word at address 2 if this is a little endian system?

0% **A.** 00110011 11011101

0% **B.** 11011101 00110011

0% **C.** 11011101 00101100

0% **D.** 00101100 11011101

0	00001111
1	00110011
2	11011101
3	00101100

✓

Handwritten note: 00101100 11011101 ✓

# What makes an ISA?

- ✓ 1. Memory models
- ✓ 2. Registers
3. Instructions
4. Data types



# What makes an ISA?

## 2: Registers

- Two types of register
  - General purpose
    - Used for temporary results etc
  - Special purpose, e.g.
    - Program Counter (PC)
    - Stack pointer (SP)
    - Input/Output Registers
    - Status Register
      - (Tanenbaum calls this Program Status Word)

## 2: Registers (cont.)

- Some other registers are part of the microarchitecture NOT the ISA
  - e.g. Instruction Register (IR) ✕
  - i.e. programmer doesn't need to know about these (and can't directly change or use them)

✓ F — PC  
✓ D  
E

# AVR Registers

- General purpose registers (0-31) are quite regular and can be used interchangeably
  - Exception: a few instructions work on half the registers (registers 16-31)
    - (Only 4 bits used to represent operand in instruction)
  - X, Y, Z registers
    - Next slide
- There are many I/O registers
  - Not to be confused with general purpose registers
  - Some instructions work with these, others with general purpose registers – don't confuse them

224


R0 - R31      8 bits  
32 reg.

# AVR X,Y,Z registers

D FFs

8 bit  
reg

R0 ✓	0x00
R1 ✓	0x01
R2 ✓	0x02
...	
R13	0x0D
R14	0x0E
R15	0x0F
R16	0x10
R17	0x11
← 8 bits ... →	
R26 ✓ { X	0x1A
R27 ✓ { X	0x1B
R28 { Y	0x1C
R29 { Y	0x1D
R30 { Z	0x1E
R31 ✓ { Z	0x1F



✓ ADD R2, R3

X-register Low Byte { X  
 X-register High Byte { X  
 Y-register Low Byte { Y  
 Y-register High Byte { Y  
 Z-register Low Byte { Z  
 Z-register High Byte { Z

16 bit reg.

# AVR I/O Registers

- AVR ATmega324A has 224 I/O register addresses to control peripherals and get data to/from them, e.g.

- ✓ ■ Timers and counters
- ✓ ■ Analog to Digital Converters
- Serial input/output
- General purpose input/output ports


- ✓ ■ Three registers associated with each (learning lab 11)

GPI/O

- ✓ ■ **DDRx** - Data direction register ✓
- ✓ ■ **PORTx** - Values to output ✓
- ✓ ■ **PINx** - Values on the pins ✓

(x = A,B,C or D)

Data Mem



- Some addresses reserved (not actually available on this device)

# What makes an ISA?

## 3: Instructions

- This is the main feature of an ISA
- Instruction types include
  - **Input/Output** – communicate with I/O devices
  - **Load/Store** – move data from/to memory
  - ✓ **Move** – copy data between registers
  - ✓ **Arithmetic** – addition, subtraction, ...
  - **Logical** – Boolean operations
  - **Branching** – for deciding which instruction to perform next

high level

✗ assembly lang

m/c code

✓ 131 ✓ assembly instruction.


# Recall - AVR Instructions

- **MOV rd, rr**
  - “move” (but actually means copy) contents of register rr (source) to register rd (destination)
  - Example: `mov r3, r14`
- **ADD rd, rr**
  - Add contents of register rr to register rd
  - Example: `add r5, r6` ;  $r5 \leftarrow r5 + r6$
- **AND rd, rr; OR rd, rr; EOR rd, rr**
  - Bitwise operations on given two registers, result goes into rd
- **LDI rd, K** (where K is a constant)
  - Load constant value into a register (16 to 31)

r16-r31

`ldi r0, 53` X  
 $r(16-r31)$

# Input/Output Instructions

- AVR ATmega324A has 224 I/O registers to control peripherals
- I/O Registers 0 to 223
  - Available at memory addresses 32 to 255 (\$20 to \$FF)
    - Use load and store instructions to access
- First 64 I/O registers (0 to 63) are special 
  - Can also be accessed using "in" and "out" instructions
- Other 160 I/O registers (64 to 223) are called "extended" I/O registers
  - Can only be accessed using load and store instructions



# IN and OUT instructions



## • **in rd, P**

- Load I/O register value into general purpose register
- rd: r0 to r31
- P = I/O register number (0 to 63, 0 to 0x3F)
- Example:

I/O reg no.


in r3, 0



## • **out P, rr**

- Store value from general purpose register into I/O register
- rr: r0 to r31, P = 0 to 63 (0 to 0x3F)
- Example:

out 1, r3



# Summary:

## What makes an ISA?

1. Memory models
  2. Registers
  3. Instructions (more in lecture 13)
  4. Data types (lecture 13)
- If you know all these details, you can
    - Write machine code that runs on the CPU
    - Build the CPU